

1. Computing:

$$\begin{aligned}
 P(\text{Target in Cell}_i | \text{Observations}_t \wedge \text{Failure in Cell}_j) &= B_{t+1} \\
 &= \frac{P(\text{Target in Cell}_i | \text{Observations}_t) P(\text{Target in Cell}_i | \text{Failure in Cell}_j)}{P(\text{Observations}_t \wedge \text{failed in } j)} \\
 &= B_t(i) * \frac{P(\text{Failure in } j | \text{Target in } i)}{P(\text{Observations}_t \wedge \text{failed in } j)}
 \end{aligned}$$

Since  $\frac{1}{P(\text{Observations}_t \wedge \text{failed in } j)}$  is just the normalization factor, it can be showed as

$\propto$ . Therefore, we can use  $B_{t+1}(i) = \propto B_t(i) P(\text{failure in } j | \text{target is in } i)$ .

Using this we can efficiently update the belief state since, we're using our prior belief and our current result for the update. Then, normalizing the belief.

2. The probability that the target would be found in Cell  $i$  if it is searched is the prior probability multiplied by the false negative rate

$\rightarrow P(\text{Target not found in Cell}_i | \text{Target in Cell}_i)$  of the current cell. Then dividing by the sum of all the evidence collected so far.

In the program, this can be done by

multiplying the current belief at the current cell by the probability of its terrain.

Then normalizing by multiplying by  $\frac{1}{\sum \text{evidence collected}}$

3. Rule 2 on average requires less searches. I think this is because rule 2 skips cells that have a low probability of finding the target within it.

That is  $P(\text{Target not in found in cell } i | \text{Target in cell } i) < 0.2$ .

This means that cells that the only cells that are searched are those that have a probability of finding the target in the cell given that it is in there  $= 0.8$ .

This holds across multiple maps. It consistently does less searches than rule 1.

4. You can use your current belief state and your current location to determine where to search next is by checking the probability of the surrounding cells containing the target and checking the probability of finding the target within that cell. This is a combination of rules 1 and 2. This means that for any Cell ( $i$ ) at most it will conduct four searches plus one action to move. The performance compared to the first two rules worse due to there being more constraints on where to move and where to search. This means that it fails to find the target more often than the other two rules, but the times that it does find the target, the amount of actions are less than the first rule.

5. This project is like the old joke by that the rules we're searching each cell by is like the light. Searching blindly in any given cell because we think that the target might be there. The light in this case tells us that even if the search is easier, the false positive rate will always be 0. We won't ever find the target, were it is not located. The "park" in this project is the terrain and the probability of finding the target within the given terrain. Some of the rules used to search was like the man, ignoring the park because it was more difficult to search there. In this case, some of the rules ignored certain terrains because the probability of finding the target in those terrains was very low, compared to others. So, even if the target was in one of the terrains that was ignored, it would never be found since we weren't even trying to search in there.